

# Student Performance Indicator

## Part 1: Life cycle of Machine learning Project

- Understanding the Problem Statement
- Data Collection
- Data Checks to perform
- Exploratory data analysis
- Data Pre-Processing
- Model Training
- Choose best model

### 1) Problem statement

- This project understands how the student's performance (test scores) is affected by other variables such as Gender, Ethnicity, Parental level of education, Lunch and Test preparation course.

### 2) Data Collection

- Dataset Source - <https://www.kaggle.com/datasets/spscientist/students-performance-in-exams?datasetId=74977>
- The data consists of 8 column and 1000 rows.

### 2.1 Import Data and Required Packages

Importing Pandas, Numpy, Matplotlib, Seaborn and Warings Library.

```
In [7]: pip install catboost
```

Requirement already satisfied: catboost in c:\users\ravi\anaconda3\lib\site-packages (1.2.2)

Requirement already satisfied: graphviz in c:\users\ravi\anaconda3\lib\site-packages (from catboost) (0.20.1)

Requirement already satisfied: matplotlib in c:\users\ravi\anaconda3\lib\site-packages (from catboost) (3.7.1)

Requirement already satisfied: numpy>=1.16.0 in c:\users\ravi\anaconda3\lib\site-packages (from catboost) (1.24.3)

Requirement already satisfied: pandas>=0.24 in c:\users\ravi\anaconda3\lib\site-packages (from catboost) (1.5.3)

Requirement already satisfied: scipy in c:\users\ravi\anaconda3\lib\site-packages (from catboost) (1.10.1)

Requirement already satisfied: plotly in c:\users\ravi\anaconda3\lib\site-packages (from catboost) (5.9.0)

Requirement already satisfied: six in c:\users\ravi\anaconda3\lib\site-packages (from catboost) (1.16.0)

Requirement already satisfied: python-dateutil>=2.8.1 in c:\users\ravi\anaconda3\lib\site-packages (from pandas>=0.24->catboost) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\ravi\anaconda3\lib\site-packages (from pandas>=0.24->catboost) (2022.7)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\ravi\anaconda3\lib\site-packages (from matplotlib->catboost) (1.0.5)

Requirement already satisfied: cyclor>=0.10 in c:\users\ravi\anaconda3\lib\site-packages (from matplotlib->catboost) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\ravi\anaconda3\lib\site-packages (from matplotlib->catboost) (4.25.0)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\ravi\anaconda3\lib\site-packages (from matplotlib->catboost) (1.4.4)

Requirement already satisfied: packaging>=20.0 in c:\users\ravi\anaconda3\lib\site-packages (from matplotlib->catboost) (23.0)

Requirement already satisfied: pillow>=6.2.0 in c:\users\ravi\anaconda3\lib\site-packages (from matplotlib->catboost) (9.4.0)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\ravi\anaconda3\lib\site-packages (from matplotlib->catboost) (3.0.9)

Requirement already satisfied: tenacity>=6.2.0 in c:\users\ravi\anaconda3\lib\site-packages (from plotly->catboost) (8.2.2)

Note: you may need to restart the kernel to use updated packages.

In [58]: `pip install xgboost`

Requirement already satisfied: xgboost in c:\users\ravi\anaconda3\lib\site-packages (2.0.3)Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: numpy in c:\users\ravi\anaconda3\lib\site-packages (from xgboost) (1.24.3)

Requirement already satisfied: scipy in c:\users\ravi\anaconda3\lib\site-packages (from xgboost) (1.10.1)

In [59]: `import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
# for modelling
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor
from sklearn.svm import SVR
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
from sklearn.model_selection import RandomizedSearchCV`

```
from catboost import CatBoostRegressor
from xgboost import XGBRegressor
import warnings
```

## Import the CSV Data as Pandas DataFrame

```
In [9]: df = pd.read_csv('stud.csv')
```

## Show Top 5 Records

```
In [11]: df.head()
```

```
Out[11]:
```

	gender	race_ethnicity	parental_level_of_education	lunch	test_preparation_course	math_s
0	female	group B	bachelor's degree	standard	none	
1	female	group C	some college	standard	completed	
2	female	group B	master's degree	standard	none	
3	male	group A	associate's degree	free/reduced	none	
4	male	group C	some college	standard	none	

## Shape of the dataset

```
In [12]: df.shape
```

```
Out[12]: (1000, 8)
```

## 2.2 Dataset information

- gender : sex of students -> (Male/female)
- race/ethnicity : ethnicity of students -> (Group A, B,C, D,E)
- parental level of education : parents' final education ->(bachelor's degree,some college,master's degree,associate's degree,high school)
- lunch : having lunch before test (standard or free/reduced)
- test preparation course : complete or not complete before test
- math score
- reading score
- writing score

## 3. Data Checks to perform

- Check Missing values
- Check Duplicates
- Check data type
- Check the number of unique values of each column
- Check statistics of data set
- Check various categories present in the different categorical column

### 3.1 Check Missing values

```
In [13]: df.isna().sum()
```

```
Out[13]: gender                0
race_ethnicity              0
parental_level_of_education  0
lunch                      0
test_preparation_course     0
math_score                  0
reading_score               0
writing_score               0
dtype: int64
```

There are no missing values in the data set

### 3.2 Check Duplicates

```
In [14]: df.duplicated().sum()
```

```
Out[14]: 0
```

There are no duplicates values in the data set

### 3.3 Check data types

```
In [15]: # Check Null and Dtypes
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   gender                                1000 non-null   object
 1   race_ethnicity                       1000 non-null   object
 2   parental_level_of_education          1000 non-null   object
 3   lunch                                1000 non-null   object
 4   test_preparation_course              1000 non-null   object
 5   math_score                           1000 non-null   int64
 6   reading_score                        1000 non-null   int64
 7   writing_score                         1000 non-null   int64
dtypes: int64(3), object(5)
memory usage: 62.6+ KB
```

### 3.4 Checking the number of unique values of each column

```
In [16]: df.nunique()
```

```
Out[16]: gender                2
race_ethnicity              5
parental_level_of_education  6
lunch                      2
test_preparation_course     2
math_score                  81
reading_score               72
writing_score               77
dtype: int64
```

### 3.5 Check statistics of data set

In [17]: `df.describe()`

Out[17]:

	math_score	reading_score	writing_score
<b>count</b>	1000.00000	1000.000000	1000.000000
<b>mean</b>	66.08900	69.169000	68.054000
<b>std</b>	15.16308	14.600192	15.195657
<b>min</b>	0.00000	17.000000	10.000000
<b>25%</b>	57.00000	59.000000	57.750000
<b>50%</b>	66.00000	70.000000	69.000000
<b>75%</b>	77.00000	79.000000	79.000000
<b>max</b>	100.00000	100.000000	100.000000

#### Insight

- From above description of numerical data, all means are very close to each other - between 66 and 68.05;
- All standard deviations are also close - between 14.6 and 15.19;
- While there is a minimum score 0 for math, for writing minimum is much higher = 10 and for reading myet higher = 17

### 3.7 Exploring Data

In [18]: `df.head()`

Out[18]:

	gender	race_ethnicity	parental_level_of_education	lunch	test_preparation_course	math_s
<b>0</b>	female	group B	bachelor's degree	standard	none	
<b>1</b>	female	group C	some college	standard	completed	
<b>2</b>	female	group B	master's degree	standard	none	
<b>3</b>	male	group A	associate's degree	free/reduced	none	
<b>4</b>	male	group C	some college	standard	none	

In [19]:

```

print("Categories in 'gender' variable:      ",end=" " )
print(df['gender'].unique())

print("Categories in 'race_ethnicity' variable: ",end=" ")
print(df['race_ethnicity'].unique())

print("Categories in 'parental level of education' variable:",end=" " )
print(df['parental_level_of_education'].unique())

print("Categories in 'lunch' variable:      ",end=" " )
print(df['lunch'].unique())

```

```
print("Categories in 'test preparation course' variable: ",end=" ")
print(df['test_preparation_course'].unique())
```

```
Categories in 'gender' variable:      ['female' 'male']
Categories in 'race_ethnicity' variable:  ['group B' 'group C' 'group A' 'group
D' 'group E']
Categories in 'parental level of education' variable: ["bachelor's degree" 'some co
llege' "master's degree" "associate's degree"
'high school' 'some high school']
Categories in 'lunch' variable:      ['standard' 'free/reduced']
Categories in 'test preparation course' variable:      ['none' 'completed']
```

```
In [20]: # define numerical & categorical columns
numeric_features = [feature for feature in df.columns if df[feature].dtype != 'O']
categorical_features = [feature for feature in df.columns if df[feature].dtype == 'O']

# print columns
print('We have {} numerical features : {}'.format(len(numeric_features), numeric_features))
print('\nWe have {} categorical features : {}'.format(len(categorical_features), categorical_features))
```

```
We have 3 numerical features : ['math_score', 'reading_score', 'writing_score']
```

```
We have 5 categorical features : ['gender', 'race_ethnicity', 'parental_level_of_education', 'lunch', 'test_preparation_course']
```

```
In [21]: df.head(2)
```

```
Out[21]:
```

	gender	race_ethnicity	parental_level_of_education	lunch	test_preparation_course	math_score
0	female	group B	bachelor's degree	standard	none	7
1	female	group C	some college	standard	completed	6

### 3.8 Adding columns for "Total Score" and "Average"

```
In [22]: df['total_score'] = df['math_score'] + df['reading_score'] + df['writing_score']
df['average'] = df['total_score']/3
df.head()
```

```
Out[22]:
```

	gender	race_ethnicity	parental_level_of_education	lunch	test_preparation_course	math_score	reading_score	writing_score	total_score	average
0	female	group B	bachelor's degree	standard	none	7	8	9	24	8
1	female	group C	some college	standard	completed	6	7	8	21	7
2	female	group B	master's degree	standard	none	8	9	10	27	9
3	male	group A	associate's degree	free/reduced	none	5	6	7	18	6
4	male	group C	some college	standard	none	6	7	8	21	7

```
In [23]: reading_full = df[df['reading_score'] == 100]['average'].count()
writing_full = df[df['writing_score'] == 100]['average'].count()
math_full = df[df['math_score'] == 100]['average'].count()

print(f'Number of students with full marks in Maths: {math_full}')
print(f'Number of students with full marks in Writing: {writing_full}')
print(f'Number of students with full marks in Reading: {reading_full}')
```

Number of students with full marks in Maths: 7  
 Number of students with full marks in Writing: 14  
 Number of students with full marks in Reading: 17

```
In [24]: reading_less_20 = df[df['reading_score'] <= 20]['average'].count()
writing_less_20 = df[df['writing_score'] <= 20]['average'].count()
math_less_20 = df[df['math_score'] <= 20]['average'].count()

print(f'Number of students with less than 20 marks in Maths: {math_less_20}')
print(f'Number of students with less than 20 marks in Writing: {writing_less_20}')
print(f'Number of students with less than 20 marks in Reading: {reading_less_20}')
```

Number of students with less than 20 marks in Maths: 4  
 Number of students with less than 20 marks in Writing: 3  
 Number of students with less than 20 marks in Reading: 1

### Insights

- From above values we get students have performed the worst in Maths
- Best performance is in reading section

## 4. Exploring Data ( Visualization )


### 4.1 Visualize average score distribution to make some conclusion.

- Histogram
- Kernel Distribution Function (KDE)

#### 4.1.1 Histogram & KDE

```
In [25]: fig, axs = plt.subplots(1, 2, figsize=(15, 7))
plt.subplot(121)
sns.histplot(data=df, x='average', bins=30, kde=True, color='g')
plt.subplot(122)
sns.histplot(data=df, x='average', kde=True, hue='gender')
plt.suptitle('Histograms and KDEs of Average Values', fontsize=16)
plt.show()
```


Histograms and KDEs of Average Values



```
In [26]: fig, axs = plt.subplots(1, 2, figsize=(15, 7))
plt.subplot(121)
```

```
sns.histplot(data=df,x='total score',bins=30,kde=True,color='g')
plt.subplot(122)
sns.histplot(data=df,x='total score',kde=True,hue='gender')
plt.suptitle('Histograms and KDEs of Total scores', fontsize=16)
plt.show()
```

Histograms and KDEs of Total scores




### Insights

- Female students tend to perform well then male students.

```
In [27]: plt.subplots(1,3,figsize=(25,6))
plt.subplot(141)
sns.histplot(data=df,x='average',kde=True,hue='lunch')
plt.subplot(142)
sns.histplot(data=df[df.gender=='female'],x='average',kde=True,hue='lunch')
plt.subplot(143)
sns.histplot(data=df[df.gender=='male'],x='average',kde=True,hue='lunch')
plt.suptitle('Histograms and KDEs of male and female performance', fontsize=16)
plt.show()
```

Histograms and KDEs of male and female performance




### Insights

- Standard lunch helps perform well in exams.
- Standard lunch helps perform well in exams be it a male or a female.



```
In [28]: plt.subplots(1,3,figsize=(25,6))
plt.subplot(141)
ax =sns.histplot(data=df,x='average',kde=True,hue='parental_level_of_education')
plt.subplot(142)
ax =sns.histplot(data=df[df.gender=='male'],x='average',kde=True,hue='parental_level_of_education')
plt.subplot(143)
ax =sns.histplot(data=df[df.gender=='female'],x='average',kde=True,hue='parental_level_of_education')
plt.suptitle('Histograms and KDEs of parental level of education male and female wise')
plt.show()
```

Histograms and KDEs of parental level of education male and female wise



### Insights


- In general parent's education don't help student perform well in exam.
- 2nd plot shows that parent's whose education is of associate's degree or master's degree their male child tend to perform well in exam
- 3rd plot we can see there is no effect of parent's education on female students.

```
In [29]: df.columns
```

```
Out[29]: Index(['gender', 'race_ethnicity', 'parental_level_of_education', 'lunch',
        'test_preparation_course', 'math_score', 'reading_score',
        'writing_score', 'total score', 'average'],
        dtype='object')
```

```
In [26]: plt.subplots(1,3,figsize=(25,6))
plt.subplot(141)
ax =sns.histplot(data=df,x='average',kde=True,hue='race_ethnicity')
plt.subplot(142)
ax =sns.histplot(data=df[df.gender=='female'],x='average',kde=True,hue='race_ethnicity')
plt.subplot(143)
ax =sns.histplot(data=df[df.gender=='male'],x='average',kde=True,hue='race_ethnicity')
plt.suptitle('Histograms and KDEs of race_ethnicity male and female wise', fontsize=14)
plt.show()
```

Histograms and KDEs of race\_ethnicity male and female wise



### Insights

- Students of group A and group B tends to perform poorly in exam.
- Students of group A and group B tends to perform poorly in exam irrespective of whether they are male or female

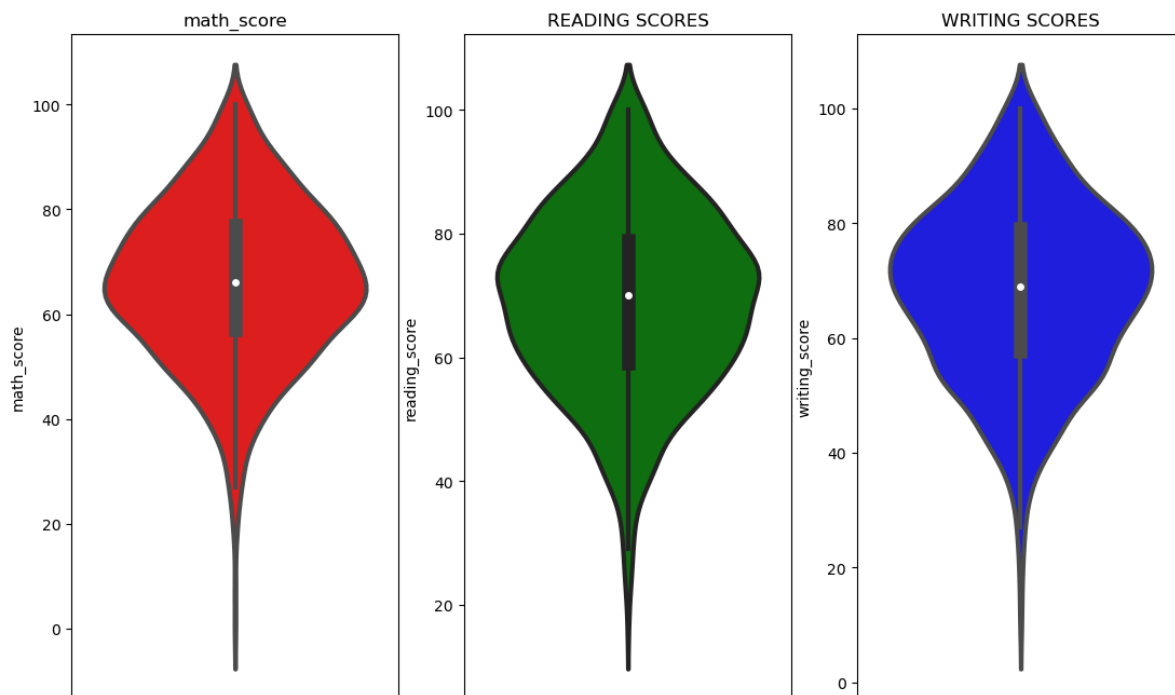
## 4.2 Maximum score of students in all three subjects

In [27]: `df.columns`

Out[27]: Index(['gender', 'race\_ethnicity', 'parental\_level\_of\_education', 'lunch', 'test\_preparation\_course', 'math\_score', 'reading\_score', 'writing\_score', 'total score', 'average'], dtype='object')

```
In [30]: plt.figure(figsize=(18,8))
plt.subplot(1, 4, 1)
plt.title('math_score')
sns.violinplot(y='math_score',data=df,color='red',linewidth=3)
plt.subplot(1, 4, 2)
plt.title('READING SCORES')
sns.violinplot(y='reading_score',data=df,color='green',linewidth=3)
plt.subplot(1, 4, 3)
plt.title('WRITING SCORES')
sns.violinplot(y='writing_score',data=df,color='blue',linewidth=3)
plt.suptitle('violin plot of math,reading and writing score', fontsize=16)
plt.show()
```

violin plot of math,reading and writing score



## Insights

- From the above three plots it's clearly visible that most of the students score in between 60-80 in Maths whereas in reading and writing most of them score from 50-80

## 4.3 Multivariate analysis using pieplot

```
In [31]: plt.rcParams['figure.figsize'] = (30, 12)

plt.subplot(1, 5, 1)
size = df['gender'].value_counts()
labels = 'Female', 'Male'
color = ['red', 'green']

plt.pie(size, colors = color, labels = labels, autopct = '%2f%')
plt.title('Gender', fontsize = 20)
plt.axis('off')

plt.subplot(1, 5, 2)
size = df['race_ethnicity'].value_counts()
labels = 'Group C', 'Group D', 'Group B', 'Group E', 'Group A'
color = ['red', 'green', 'blue', 'cyan', 'orange']

plt.pie(size, colors = color, labels = labels, autopct = '%2f%')
plt.title('Race_Ethnicity', fontsize = 20)
plt.axis('off')

plt.subplot(1, 5, 3)
size = df['lunch'].value_counts()
labels = 'Standard', 'Free'
color = ['red', 'green']
```

```
plt.pie(size, colors = color, labels = labels, autopct = '%.2f%%')
plt.title('Lunch', fontsize = 20)
plt.axis('off')

plt.subplot(1, 5, 4)
size = df['test_preparation_course'].value_counts()
labels = 'None', 'Completed'
color = ['red', 'green']


plt.pie(size, colors = color, labels = labels, autopct = '%.2f%%')
plt.title('Test Course', fontsize = 20)
plt.axis('off')

plt.subplot(1, 5, 5)
size = df['parental_level_of_education'].value_counts()
labels = 'Some College', "Associate's Degree", 'High School', 'Some High School', "Bachelor's Degree"
color = ['red', 'green', 'blue', 'cyan', 'orange', 'grey']

plt.pie(size, colors = color, labels = labels, autopct = '%.2f%%')
plt.title('Parental Education', fontsize = 20)
plt.axis('off')

plt.tight_layout()
plt.grid()

plt.show()
```



## Insights

- Number of Male and Female students is almost equal
- Number students are greatest in Group C
- Number of students who have standard lunch are greater
- Number of students who have not enrolled in any test preparation course is greater
- Number of students whose parental education is "Some College" is greater followed closely by "Associate's Degree"

## 4.4 Feature Wise Visualization

### 4.4.1 GENDER COLUMN

- How is distribution of Gender ?
- Is gender has any impact on student's performance ?


## UNIVARIATE ANALYSIS ( How is distribution of Gender ? )

```
In [32]: f,ax=plt.subplots(1,2,figsize=(20,10))
sns.countplot(x=df['gender'],data=df,palette = 'bright',ax=ax[0],saturation=0.95)
for container in ax[0].containers:
```

```
ax[0].bar_label(container,color='black',size=20)
plt.suptitle('How is distribution of Gender', fontsize=16)

plt.pie(x=df['gender'].value_counts(),labels=['Male','Female'],explode=[0,0.1],auto
plt.show()
```

How is distribution of Gender



## Insights

- Gender has balanced data with female students are 518 (48%) and male students are 482 (52%)

## BIVARIATE ANALYSIS ( Is gender has any impact on student's performance ? )

```
In [33]: gender_group = df.groupby('gender').mean()
gender_group
```

```
Out[33]:
```

	math_score	reading_score	writing_score	total score	average
female	63.633205	72.608108	72.467181	208.708494	69.569498
male	68.728216	65.473029	63.311203	197.512448	65.837483

```
In [34]: plt.figure(figsize=(10, 8))

X = ['Total Average', 'Math Average']

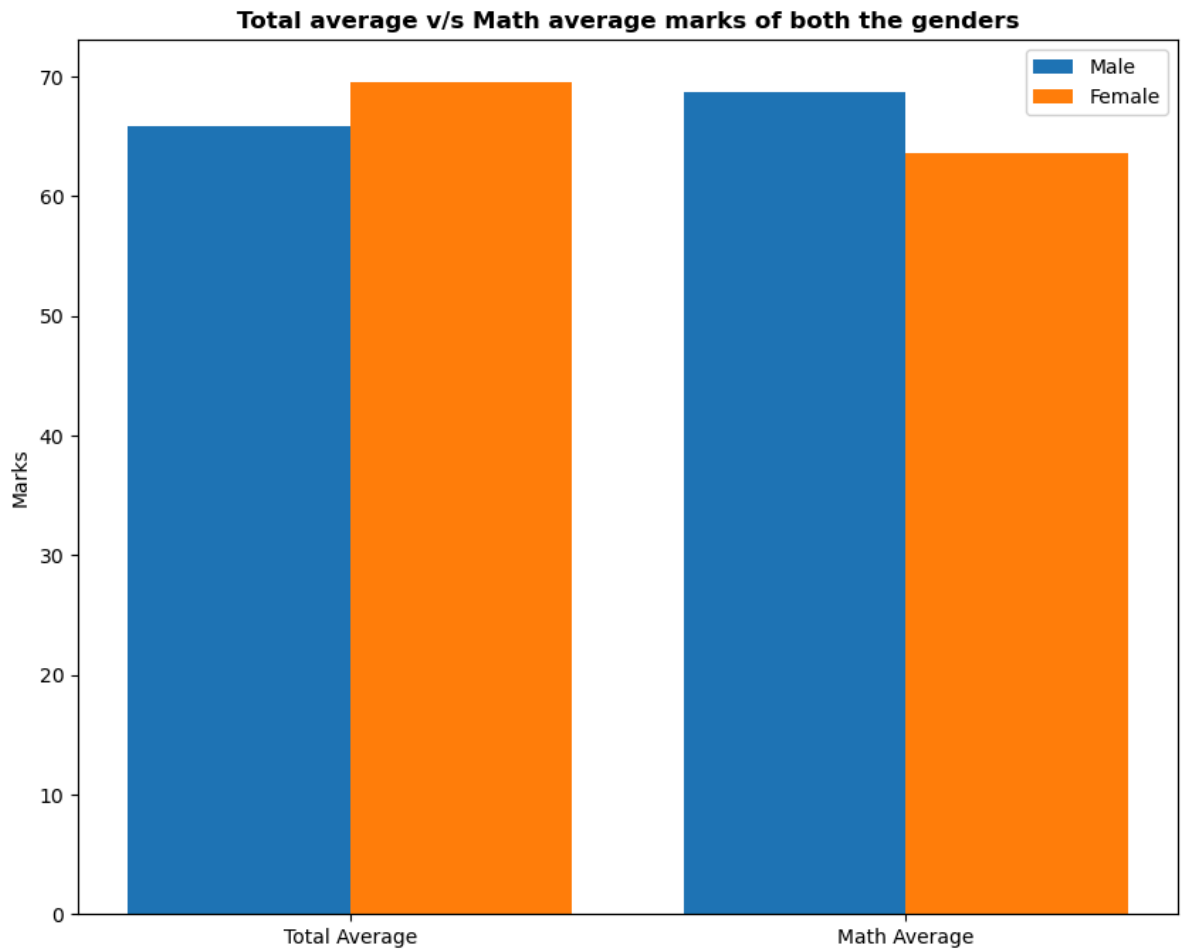
female_scores = [gender_group['average'][0], gender_group['math_score'][0]]
male_scores = [gender_group['average'][1], gender_group['math_score'][1]]

X_axis = np.arange(len(X))

plt.bar(X_axis - 0.2, male_scores, 0.4, label = 'Male')
plt.bar(X_axis + 0.2, female_scores, 0.4, label = 'Female')

plt.xticks(X_axis, X)
plt.ylabel("Marks")
```

```
plt.title("Total average v/s Math average marks of both the genders", fontweight='b')
plt.legend()
plt.show()
```



## Insights

- On an average females have a better overall score than men.
- whereas males have scored higher in Maths.



## 4.4.2 RACE/EHNICITY COLUMN

- How is Group wise distribution ?
- Is Race/Ehnicity has any impact on student's performance ?

## UNIVARIATE ANALYSIS ( How is Group wise distribution ?)

```
In [36]: f,ax=plt.subplots(1,2,figsize=(20,10))
sns.countplot(x=df['race_ethnicity'],data=df,palette = 'bright',ax=ax[0],saturation=0.8)
for container in ax[0].containers:
    ax[0].bar_label(container,color='black',size=20)
plt.suptitle('How is Group wise distribution', fontsize=16)
plt.pie(x = df['race_ethnicity'].value_counts(),labels=df['race_ethnicity'].value_counts().index)
plt.show()
```

How is Group wise distribution



## Insights

- Most of the student belonging from group C /group D.
- Lowest number of students belong to group A.

## BIVARIATE ANALYSIS ( Is Race/Ehnicity has any impact on student's performance ? )

```
In [37]: Group_data2=df.groupby('race_ethnicity')
f,ax=plt.subplots(1,3,figsize=(20,8))
sns.barplot(x=Group_data2['math_score'].mean().index,y=Group_data2['math_score'].me
ax[0].set_title('Math score',color='#005ce6',size=20)

for container in ax[0].containers:
    ax[0].bar_label(container,color='black',size=15)

sns.barplot(x=Group_data2['reading_score'].mean().index,y=Group_data2['reading_scor
ax[1].set_title('Reading score',color='#005ce6',size=20)

for container in ax[1].containers:
    ax[1].bar_label(container,color='black',size=15)

sns.barplot(x=Group_data2['writing_score'].mean().index,y=Group_data2['writing_scor
ax[2].set_title('Writing score',color='#005ce6',size=20)

plt.suptitle("Is Race/Ehnicity has any impact on student's performance", fontsize=1

for container in ax[2].containers:
    ax[2].bar_label(container,color='black',size=15)
```

Is Race/Ethnicity has any impact on student's performance



## Insights

- Group E students have scored the highest marks.
- Group A students have scored the lowest marks.
- Students from a lower Socioeconomic status have a lower avg in all course subjects

## 4.4.3 PARENTAL LEVEL OF EDUCATION COLUMN

- What is educational background of student's parent ?
- Is parental education has any impact on student's performance ?

## UNIVARIATE ANALYSIS ( What is educational background of student's parent ? )

```
In [38]: df.columns
```

```
Out[38]: Index(['gender', 'race_ethnicity', 'parental_level_of_education', 'lunch',
        'test_preparation_course', 'math_score', 'reading_score',
        'writing_score', 'total score', 'average'],
        dtype='object')
```

## Insights


- Largest number of parents are from some college.

## BIVARIATE ANALYSIS ( Is parental education has any impact on student's performance ? )

```
In [39]: df.groupby('parental_level_of_education').agg('mean').plot(kind='barh',figsize=(10,
        plt.suptitle("Is parental education has any impact on student's performance", fonts
        plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
        plt.show())
```



## Is parental education has any impact on student's performance



## Insights

- The score of student whose parents possess master and bachelor level education are higher than others.

## 4.4.4 LUNCH COLUMN

- Which type of lunch is most common among students ?
- What is the effect of lunch type on test results?

## UNIVARIATE ANALYSIS ( Which type of lunch is most common among students ? )

## Insights

- Students being served Standard lunch was more than free lunch


## BIVARIATE ANALYSIS ( Is lunch type intake has any impact on student's performance ? )

```
In [37]: f,ax=plt.subplots(1,2,figsize=(20,8))
sns.countplot(x=df['parental_level_of_education'],data=df,palette = 'bright',hue='t
ax[0].set_title('Students vs test preparation course ',color='black',size=25)
for container in ax[0].containers:
    ax[0].bar_label(container,color='black',size=20)

sns.countplot(x=df['parental_level_of_education'],data=df,palette = 'bright',hue='l
for container in ax[1].containers:
```

```
ax[1].bar_label(container,color='black',size=20)
plt.suptitle("Is lunch type intake has any impact on student's performance", fontsize=16)
```

Out[37]: Text(0.5, 0.98, "Is lunch type intake has any impact on student's performance")



## Insights

- Students who get Standard Lunch tend to perform better than students who got free/reduced lunch

## 4.4.5 TEST PREPARATION COURSE COLUMN

- Which type of lunch is most common among students ?
- Is Test preparation course has any impact on student's performance ?

## BIVARIATE ANALYSIS ( Is Test preparation course has any impact on student's performance ? )

```
In [40]: plt.figure(figsize=(18, 6)) # Adjust the figure size as needed


# Subplot 1
plt.subplot(1, 3, 1)
sns.barplot(x=df['lunch'], y=df['math_score'], hue=df['test_preparation_course'])
plt.title('Math Score')

# Subplot 2
plt.subplot(1, 3, 2)
sns.barplot(x=df['lunch'], y=df['reading_score'], hue=df['test_preparation_course'])
plt.title('Reading Score')

# Subplot 3
plt.subplot(1, 3, 3)
sns.barplot(x=df['lunch'], y=df['writing_score'], hue=df['test_preparation_course'])
plt.title('Writing Score')

plt.suptitle("Is Test Preparation Course Impacting Student's Performance", fontsize=16)
plt.show()
```

## Is Test Preparation Course Impacting Student's Performance



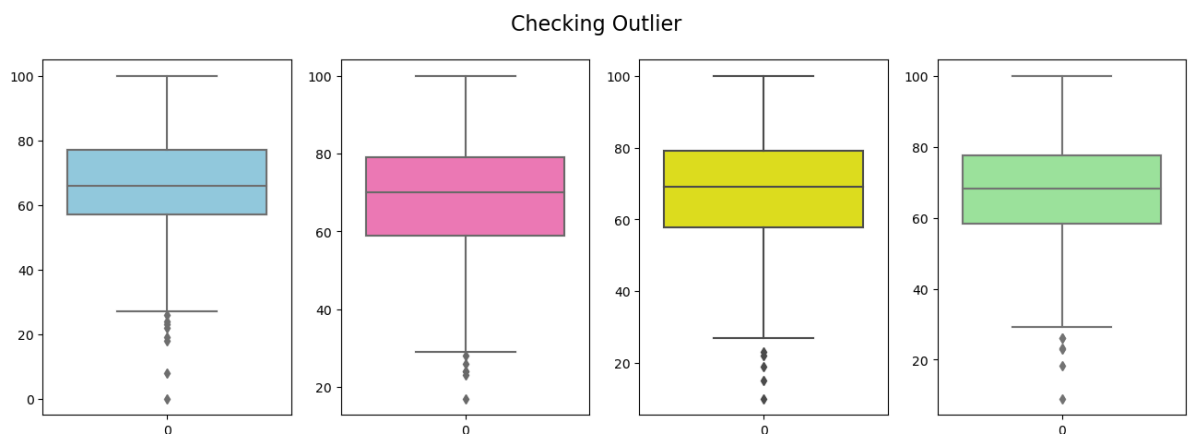
## Insights

- Students who have completed the Test Preparation Course have scores higher in all three categories than those who haven't taken the course

## 4.4.6 CHECKING OUTLIERS

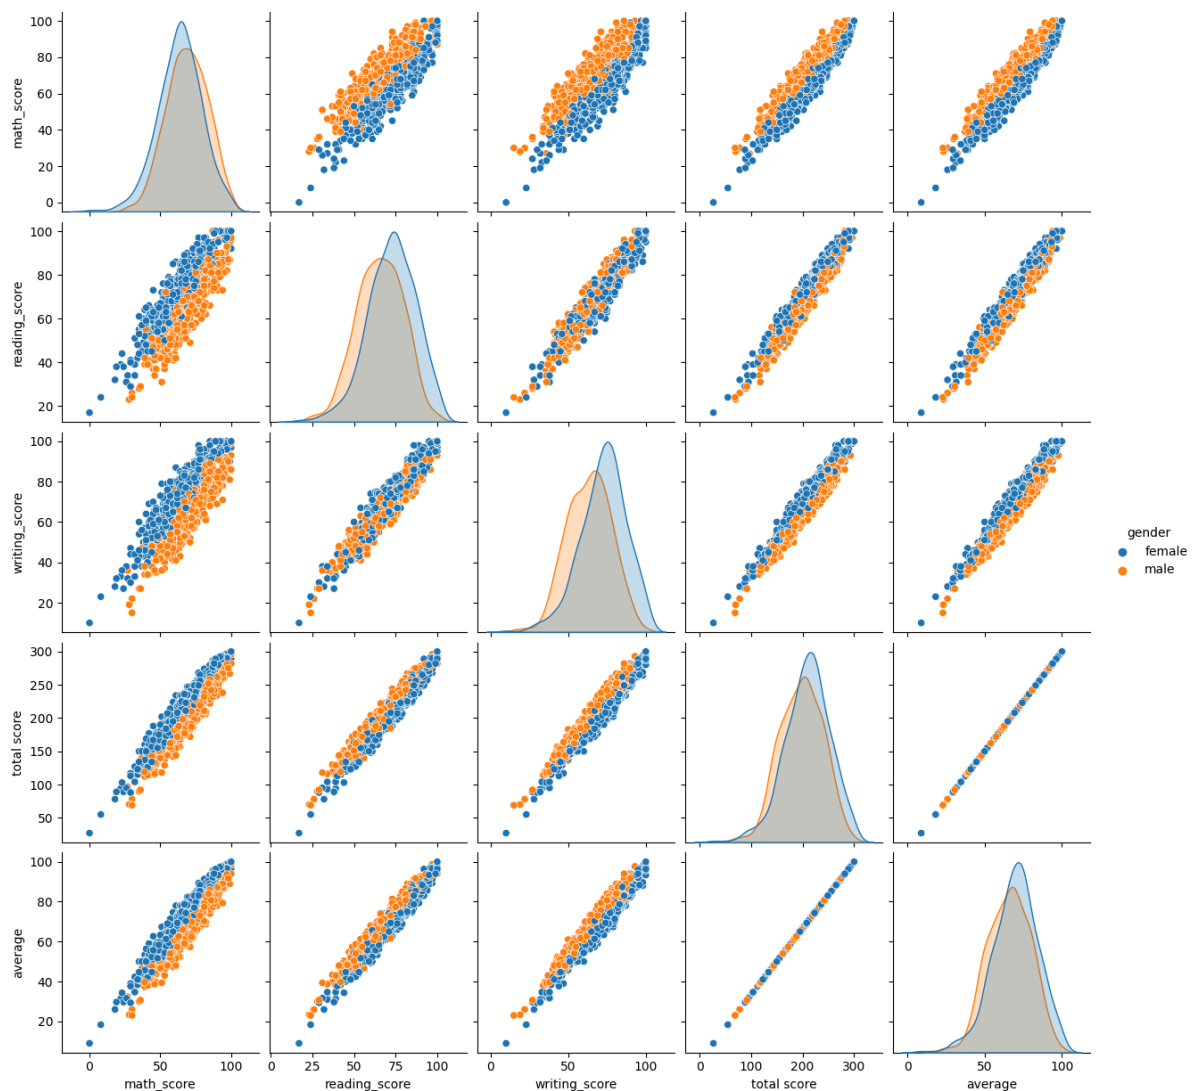
```
In [41]: plt.subplots(1,4,figsize=(16,5))
plt.subplot(141)
sns.boxplot(df['math_score'],color='skyblue')
plt.subplot(142)
sns.boxplot(df['reading_score'],color='hotpink')
plt.subplot(143)
sns.boxplot(df['writing_score'],color='yellow')
plt.subplot(144)
sns.boxplot(df['average'],color='lightgreen')

plt.suptitle("Checking Outlier", fontsize=16)
plt.show()
```



## 4.4.7 MUTIVARIATE ANALYSIS USING PAIRPLOT

```
In [40]: sns.pairplot(df,hue = 'gender')
plt.show()
```



## Insights

- From the above plot it is clear that all the scores increase linearly with each other.

## 5. Conclusions

- Student's Performance is related with lunch, race, parental level education
- Females lead in pass percentage and also are top-scorers
- Student's Performance is not much related with test preparation course
- Finishing preparation course is beneficial.

## Part 2: Model Training

```
In [42]: df = pd.read_csv('stud.csv')
```

```
In [43]: X = df.drop(columns=['math_score'], axis=1)
```

```
In [44]: X
```

Out[44]:

	gender	race_ethnicity	parental_level_of_education	lunch	test_preparation_course	read
0	female	group B	bachelor's degree	standard		none
1	female	group C	some college	standard		completed
2	female	group B	master's degree	standard		none
3	male	group A	associate's degree	free/reduced		none
4	male	group C	some college	standard		none
...	...	...	...	...		...
995	female	group E	master's degree	standard		completed
996	male	group C	high school	free/reduced		none
997	female	group C	high school	free/reduced		completed
998	female	group D	some college	standard		completed
999	female	group D	some college	free/reduced		none

1000 rows × 7 columns

In [45]: `df.head()`

Out[45]:

	gender	race_ethnicity	parental_level_of_education	lunch	test_preparation_course	math_
0	female	group B	bachelor's degree	standard		none
1	female	group C	some college	standard		completed
2	female	group B	master's degree	standard		none
3	male	group A	associate's degree	free/reduced		none
4	male	group C	some college	standard		none

In [46]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   gender                                1000 non-null   object
 1   race_ethnicity                        1000 non-null   object
 2   parental_level_of_education            1000 non-null   object
 3   lunch                                  1000 non-null   object
 4   test_preparation_course                1000 non-null   object
 5   math_score                            1000 non-null   int64
 6   reading_score                          1000 non-null   int64
 7   writing_score                           1000 non-null   int64
dtypes: int64(3), object(5)
memory usage: 62.6+ KB
```

In [47]: `y = df['math_score']`

In [48]: `categorical_features = df.select_dtypes(include='object').columns`  
`numerical_features = df.select_dtypes(exclude='object').columns`

```
print(categorical_features)
print(numerical_features)
```

```
Index(['gender', 'race_ethnicity', 'parental_level_of_education', 'lunch',
      'test_preparation_course'],
      dtype='object')
Index(['math_score', 'reading_score', 'writing_score'], dtype='object')
```

```
In [49]: num_features = X.select_dtypes(exclude='object').columns
cat_features = X.select_dtypes(include='object').columns

from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer

num_transformer = StandardScaler()
oh_transformer = OneHotEncoder()

preprocessor = ColumnTransformer(
    [
        ("OneHotEncoder", oh_transformer, cat_features),
        ("StandardScaler", num_transformer, num_features)
    ]
)
X = preprocessor.fit_transform(X)
```

```
In [50]: X.shape
```

```
Out[50]: (1000, 19)
```

```
In [51]: X
```

```
Out[51]: array([[ 1.          ,  0.          ,  0.          , ...,  1.          ,
        0.19399858,  0.39149181],
       [ 1.          ,  0.          ,  0.          , ...,  0.          ,
        1.42747598,  1.31326868],
       [ 1.          ,  0.          ,  0.          , ...,  1.          ,
        1.77010859,  1.64247471],
       ...,
       [ 1.          ,  0.          ,  0.          , ...,  0.          ,
        0.12547206, -0.20107904],
       [ 1.          ,  0.          ,  0.          , ...,  0.          ,
        0.60515772,  0.58901542],
       [ 1.          ,  0.          ,  0.          , ...,  1.          ,
        1.15336989,  1.18158627]])
```

```
In [52]: X.max(), X.min()
```

```
Out[52]: (2.112741202570347, -3.82234534162361)
```

```
In [53]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2, random_state=42)
X_train.shape,X_test.shape
```

```
Out[53]: ((800, 19), (200, 19))
```

## Evaluation function

```
In [54]: def evaluate_model(true, predicted):
mae = mean_absolute_error(true, predicted)
mse = mean_squared_error(true, predicted)
rmse = np.sqrt(mean_squared_error(true, predicted))
```

```

r2_square = r2_score(true, predicted)
return mae, rmse, r2_square

```

```

In [60]: models = {
    "Linear Regression": LinearRegression(),
    "Lasso": Lasso(),
    "Ridge": Ridge(),
    "K-Neighbors Regressor": KNeighborsRegressor(),
    "Decision Tree": DecisionTreeRegressor(),
    "Random Forest Regressor": RandomForestRegressor(),
    "XGBRegressor": XGBRegressor(),
    "CatBoosting Regressor": CatBoostRegressor(verbose=False),
    "AdaBoost Regressor": AdaBoostRegressor()
}

model_list = []
r2_list = []

for i in range(len(list(models))):
    model = list(models.values())[i]
    model.fit(X_train,y_train)

    #make prediction
    y_train_pred = model.predict(X_train)
    y_test_pred = model.predict(X_test)

    # Evaluate train and test dataset
    model_train_mae, model_train_rmse, model_train_r2 = evaluate_model(y_train, y_train_pred)
    model_test_mae, model_test_rmse, model_test_r2 = evaluate_model(y_test, y_test_pred)

    print(list(models.keys())[i])
    model_list.append(list(models.keys())[i])
    r2_list.append(round(model_test_r2,3))

    print("Model performance for Train set")
    print("- Root mean squared error : {:.4f}".format(model_train_rmse))
    print("- Mean absolute error : {:.4f}".format(model_train_mae))
    print("- R2 score : {:.4f}".format(model_train_r2))

    print("-"*35)

    print("Model performance for Test set")
    print("- Root Mean squared error: {:.4f}".format(model_test_rmse))
    print("- Mean absolute error: {:.4f}".format(model_test_mae))
    print("- R2 score : {:.4f}".format(model_test_r2))
    print('='*35)
    print("\n")

```

### Linear Regression

Model performance for Train set

- Root mean squared error : 5.3240
- Mean absolute error : 4.2691
- R2 score : 0.8743

-----  
Model performance for Test set

- Root Mean squared error: 5.3773
- Mean absolute error: 4.2053
- R2 score : 0.8812

=====

### Lasso

Model performance for Train set

- Root mean squared error : 6.5938
- Mean absolute error : 5.2063
- R2 score : 0.8071

-----  
Model performance for Test set

- Root Mean squared error: 6.5197
- Mean absolute error: 5.1579
- R2 score : 0.8253

=====

### Ridge

Model performance for Train set

- Root mean squared error : 5.3233
- Mean absolute error : 4.2650
- R2 score : 0.8743

-----  
Model performance for Test set

- Root Mean squared error: 5.3904
- Mean absolute error: 4.2111
- R2 score : 0.8806

=====

### K-Neighbors Regressor

Model performance for Train set

- Root mean squared error : 5.7133
- Mean absolute error : 4.5213
- R2 score : 0.8552

-----  
Model performance for Test set

- Root Mean squared error: 7.2488
- Mean absolute error: 5.6310
- R2 score : 0.7841

=====

### Decision Tree

Model performance for Train set

- Root mean squared error : 0.2795
- Mean absolute error : 0.0187
- R2 score : 0.9997

-----  
Model performance for Test set

- Root Mean squared error: 7.5776
- Mean absolute error: 5.9700
- R2 score : 0.7640

=====



```
Random Forest Regressor
Model performance for Train set
- Root mean squared error : 2.3223
- Mean absolute error : 1.8461
- R2 score : 0.9761
```

```
-----
Model performance for Test set
- Root Mean squared error: 5.9204
- Mean absolute error: 4.5573
- R2 score : 0.8560
=====
```

```
XGBRegressor
Model performance for Train set
- Root mean squared error : 1.0073
- Mean absolute error : 0.6875
- R2 score : 0.9955
```

```
-----
Model performance for Test set
- Root Mean squared error: 6.4733
- Mean absolute error: 5.0577
- R2 score : 0.8278
=====
```

```
CatBoosting Regressor
Model performance for Train set
- Root mean squared error : 3.0427
- Mean absolute error : 2.4054
- R2 score : 0.9589
```

```
-----
Model performance for Test set
- Root Mean squared error: 6.0086
- Mean absolute error: 4.6125
- R2 score : 0.8516
=====
```

```
AdaBoost Regressor
Model performance for Train set
- Root mean squared error : 5.9072
- Mean absolute error : 4.8231
- R2 score : 0.8452
```

```
-----
Model performance for Test set
- Root Mean squared error: 6.0436
- Mean absolute error: 4.7032
- R2 score : 0.8499
=====
```

```
In [62]: print(model_list, r2_list)
```

```
['Linear Regression', 'Lasso', 'Ridge', 'K-Neighbors Regressor', 'Decision Tree',
 'Random Forest Regressor', 'XGBRegressor', 'CatBoosting Regressor', 'AdaBoost Regr
essor'] [0.88, 0.825, 0.881, 0.783, 0.72, 0.85, 0.828, 0.852, 0.856]
```

```
In [63]: results = pd.DataFrame(list(zip(model_list, r2_list)), columns=["Model Name", 'R2_s
```

```
In [64]: results
```

Out[64]:

	Model Name	R2_Score
2	Ridge	0.881
0	Linear Regression	0.880
8	AdaBoost Regressor	0.856
7	CatBoosting Regressor	0.852
5	Random Forest Regressor	0.850
6	XGBRegressor	0.828
1	Lasso	0.825
3	K-Neighbors Regressor	0.783
4	Decision Tree	0.720


## Part 3: Final Model ( Linear Regression )

```
In [62]: lin_model = LinearRegression(fit_intercept=True)
lin_model = lin_model.fit(X_train, y_train)
y_pred = lin_model.predict(X_test)
score = r2_score(y_test, y_pred)*100
print(" Accuracy of the model is %.2f" %score)
```


Accuracy of the model is 88.12

## Plot predicted and actual value

```
In [63]: plt.scatter(y_test,y_pred);
plt.xlabel('Actual');
plt.ylabel('Predicted');
```



```
In [64]: sns.regplot(x=y_test,y=y_pred,ci=None,color='red');
```



## Difference Between Actual and Predicted value

```
In [65]: pred_df=pd.DataFrame({'Actual Value':y_test,'Predicted Value':y_pred,'Difference':y_pred_df
```

Out[65]:

	Actual Value	Predicted Value	Difference
521	91	76.507812	14.492188
737	53	58.796875	-5.796875
740	80	76.976562	3.023438
660	74	76.984375	-2.984375
411	84	87.664062	-3.664062
...	...	...	...
408	52	43.367188	8.632812
332	62	62.156250	-0.156250
208	74	67.812500	6.187500
613	65	67.125000	-2.125000
78	61	62.343750	-1.343750

200 rows × 3 columns

```
In [ ]:
```